



# Hive

A Warehousing Solution Over a Map-Reduce Framework

# Agenda

- Why Hive?
- What is Hive?
- Hive Data Model
- Hive Commands
- Hive Shell Commands
- Hive Drivers
- HiveQL
- Pros and Cons

# Challenges that Data Analysts faced


- Data Explosion
  - TBs of data generated everyday

Solution – HDFS to store data and Hadoop Map-Reduce framework to parallelize processing of Data

What is the catch?

- Hadoop Map Reduce is Java intensive
- Thinking in Map Reduce paradigm can get tricky

# Hive Key Principles



# HiveQL

DDL :

CREATE DATABASE  
CREATE TABLE  
ALTER TABLE  
SHOW TABLE  
DESCRIBE

DML:

INSERT

QUERY:

SELECT  
GROUP BY  
JOIN

# Hive Data Model

Data in Hive organized into :

- Tables
- Partitions

# Hive Data Model Contd.

- Tables

- Analogous to relational tables
- Each table has a corresponding directory in HDFS
- Data serialized and stored as files within that directory
- External Vs Internal(Managed Tables) Table

# Hive Data Types: Numeric

Type	Memory allocation
TINYINT	Its 1-byte signed integer (-128 to 127)
SMALLINT	2-byte signed integer (-32768 to 32767)
INT	4 –byte signed integer ( - 2,147,484,648 to 2,147,484,647)
BIGINT	8 byte signed integer
FLOAT	4 – byte single precision floating point number
DOUBLE	8- byte double precision floating point number
DECIMAL	We can define precision and scale in this Type



# Hive Data Types: String

Type	Length
CHAR	255
VARCHAR	1 to 65355
STRING	We can define length here(No Limit)

# Hive Data Types: Date/Date time

Type	Usage
Timestamp	Supports traditional <u>Unix</u> timestamp with optional nanosecond precision
Date	<ul style="list-style-type: none"><li>• It's in YYYY-MM-DD format.</li></ul>

> `Select current_date();` -- gives current date

> `Select current_timestamp();` -- gives current timestamp

# Hive Data Model Contd.

- Partitions

- Each table can be broken into partitions
- Partitions determine distribution of data within subdirectories

# Hive Commands

## **#list all databases available**

```
hive> show databases;
```

## **#Create a database**

```
hive> create database hadoop;
```

## **#Select a database**

```
Hive> use hadoop;
```

## **# Show list of tables**

```
hive> show tables;
```

# Hive Commands – External Table

## #create External table

```
CREATE External TABLE vbajaj.Sales (sale_id INT, amount FLOAT)  
ROW FORMAT DELIMITED  
FIELDS TERMINATED BY '\t'  
LINES TERMINATED BY '\n' ;
```

## #describe table structure

```
> desc sales;
```

## #insert Values

```
> insert into table sales select * from (select 1,111.01)a;
```

# Hive Commands – External Table

**#get hdfs location for table**

> show create table sales;

**#get hdfs files list**

> hdfs dfs -ls /user/hive/warehouse/sales/

**#cat the above file**

hdfs dfs -cat /user/hive/warehouse/sales/000000\_0

# Hive Commands- Managed Tables

## #create table

```
CREATE TABLE Sales_int (sale_id INT, amount FLOAT)
ROW FORMAT DELIMITED
FIELDS TERMINATED BY '\t'
LINES TERMINATED BY '\n'
STORED AS TEXTFILE;
```

```
> insert into table sales_int select * from (select
6,890.01000000676767576576576587586)a;
```

```
> select * from sales_int;
```

# Hive Commands – Managed Tables

## #create table

```
CREATE TABLE Sales_part (amount FLOAT)
PARTITIONED BY (sale_id INT )
ROW FORMAT DELIMITED
FIELDS TERMINATED BY '\t'
LINES TERMINATED BY '\n'
STORED AS TEXTFILE;
```

```
> set hive.exec.dynamic.partition.mode=nonstrict;
> insert into table sales_part partition(sales_id) select amt,id
from (select 1 id ,111.01 as amt)a;
\
```

So each partition will be split out into different folders like  
Sales\_part/sales\_id=1



# Hive Shell Commands

```
hive -e "show databases;"
```

```
hive -e "show tables;"
```

```
hive -e "Select * from sales;"
```

```
hive -e "insert into sales select * from (select 3,111.01)a;"
```

# Hive Shell Commands


##for running files

```
CREATE TABLE hadoop.my_table (id INT, amount decimal(38,12))  
ROW FORMAT DELIMITED  
FIELDS TERMINATED BY '\t'  
LINES TERMINATED BY '\n'  
STORED AS TEXTFILE;
```

```
insert into hadoop.my_table select * from (select 1,  
44907375.5787794758857897598)a;
```

hive -f 'create\_my\_table.sql'

# Hive Driver



- **Driver** - Maintains the lifecycle of HiveQL statement
- **Query Compiler** – Compiles HiveQL in a DAG of map reduce tasks
- **Executor** - Executes the tasks plan generated by the compiler in proper dependency order. Interacts with the underlying Hadoop instance

# Advantages

- Boon for Data Analysts
- Easy Learning curve
- Partitions(speed!)
- Flexibility to load data from localFS/HDFS into Hive Tables

# Cons and Possible Improvements

- Extending the SQL queries support(Updates, Deletes)
- Explore methods for multi query optimization

# REFERENCES

- <https://hive.apache.org/>
- <https://cwiki.apache.org/confluence/display/Hive/Presentations>
- <https://developer.yahoo.com/blogs/hadoop/comparing-pig-latin-sql-constructing-data-processing-pipelines-444.html>
- <http://www.qubole.com/blog/big-data/hive-best-practices/>